



Missing Data

Data Science & Machine Learning Team

03/31/2021

Jingjie Zhang, PhD

jingjie.zhang@hms.com

Agenda

- Why Missing Data Occurs?
 - Missing Data Mechanisms
 - Formats of Missing Data
- Handling Missing Data in Python/Pandas
- Handling Missing Data Strategies
 - Dropping Cases/Columns
 - Mean/Median/Mode Imputation
 - Predicting Missing Cases Using Machine Learning
 - Multiple Imputation Is For Inference, Not For Prediction
- Examples in Python

Why Missing Data Occurs?

Three general “missing mechanisms” moving from the simplest to the general:

- Missing Completely at Random (MCAR)
 - If the probability of missing data is the same for all units.
 - Deletion missing data does not bias your inference.
 - Roll a dice; lottery number, ICD10 - CMs ...
- Missing at Random (MAR)
 - If the probability of missing data depends on a set of observed responses.
 - Most common, missing values can be excluded (treated as NAs) or imputed.
 - CPT depends on ICD10-CM.
- Missing Not at Random (MNAR)
 - If the mechanism of missing data does not meet MCAR or MAR.
 - The only way to obtain unbiased estimates is to model the missing data process.
 - Covid-19 Symptoms

Formats of Missing Data

The presence of missing data:

- Null
 - Absence of everything; missing; empty
 - In HMS EDW_AR_FL, missing data format.
- Data-specific convention
 - Blank “” or “ ” or any invisible characters.
 - -9999 or -1
 - Boolean mask: True/False; 0/1
 - “?” In HMS EDW_CTS_FL, missing data format.
- Global convention
 - Nan (numpy nan type – np.nan, IEEE floating-point specification) or None (python object)
 - Most common

Pandas default
Missing data formats

Working with Missing Data in Numpy/Pandas

array- None

- None can only present in arrays with data type “object”

```
array1 = np.array([1, 2, None, 3, ])  
array1  
  
array([1, 2, None, 3], dtype=object)
```

- Operations on python “object” type is much slower than operations on arrays with native types

```
for dtype in ['object', 'int']:  
    print("dtype =", dtype)  
    %timeit np.arange(1E6, dtype = dtype).sum()  
    print()  
  
dtype = object  
56 ms ± 2.81 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)  
  
dtype = int  
2.11 ms ± 196 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)
```

- Cannot perform aggregation (sum/min/max/avg) across array with Nonetype

```
array1.sum()  
  
TypeError: unsupported operand type(s) for +: 'int' and 'NoneType'
```

Working with Missing Data in Numpy/Pandas

array- NaN

- NaN (Not a Number): a special floating-point value recognized by all systems

```
array2 = np.array([1, 2, np.nan, 3])
array2.dtype
dtype('float64')
```

- Result of arithmetic with NaN will be another NaN

```
print(1 + np.nan, 1*np.nan, array2.sum(), array2.min())
nan nan nan nan
```

- Nan is a floating-point value; there is no equivalent NaN for integers, strings, or other types.

None & NaN in Pandas

- They are **interchangeable**

```
df = pd.Series([1, np.nan, 2, None])
df
0    1.0
1    NaN
2    2.0
3    NaN
dtype: float64
```

- Up-casting conventions in Pandas when NA values are introduced:

Typeclass	Conversion When Storing NAs	NA Sentinel Value
object	No change	None or np.nan
float	No change	np.nan
int	cast to float	np.nan
boolean	No change	<NA>

```
for dtype in ['object', 'float', 'int', 'boolean']:
    df = pd.Series([0, 1], dtype = dtype)
    df[0] = np.nan
    print("Original Series dtype = ", dtype)
    print(df, '\n')
```

```
Original Series dtype = object
0    NaN
1     1
dtype: object
```

```
Original Series dtype = float
0    NaN
1    1.0
dtype: float64
```

```
Original Series dtype = int
0    NaN
1    1.0
dtype: float64
```

```
Original Series dtype = boolean
0    <NA>
1    True
dtype: boolean
```

Handling Missing Data in Pandas

- Detecting null values - `isnull()` / `isna()` / `notnull()` / `notna()`

df	df.isnull()	df.notna()	df[df.notna()]
0 1.0	0 False	0 True	0 1.0
1 NaN	1 True	1 False	2 2.0
2 2.0	2 False	2 True	dtype: float64
3 NaN	3 True	3 False	
dtype: float64	dtype: bool	dtype: bool	

- Dropping null values – `dropna()`

df.dropna(axis = 1)

	A	B	C
0	NaN	1.0	1
1	2.0	2.0	2
2	3.0	NaN	3

df.dropna(axis = 0)

	A	B	C
1	2.0	2.0	2

- Filling null values – `fillna()`

df	df.fillna(0)	df.fillna(method = 'bfill')	df.fillna(method = 'ffill')
0 1.0	0 1.0	0 1.0	0 1.0
1 NaN	1 0.0	1 2.0	1 1.0
2 2.0	2 2.0	2 2.0	2 2.0
3 NaN	3 0.0	3 3.0	3 2.0
4 3.0	4 3.0	4 3.0	4 3.0
dtype: float64	dtype: float64	dtype: float64	dtype: float64

Next Observation
Carried
Backward(NOCB)

Last Observation
Carried
Forward(LOCF)

Handling Strategies - Deletion

- Listwise deletion (aka complete case analysis)
 - Simply drop all rows/samples containing missing values.
 - Pros: easy to implement
 - Cons: loss of data; increase the standard error and widen the confidence interval.
- Drop columns/fields/variables
 - Simply drop columns with majority of data are missing values
 - Be cautious using this approach

Handling Strategies - Deletion

Listwise Deletion

Gender	Age	Weight
F	20	130
M	NaN	150
F	30	132
M	40	160
F	43	NaN
F	50	150

```
df.dropna(axis = 0)
```

Drop Column

Gender	Age	Weight
F	20	130
M	NaN	150
F	30	132
M	40	160
F	43	NaN
F	50	150

```
df.dropna(axis = 1)
```

Handling Strategies – Mean/Median/Mode Imputation

Replace missing values with the variable mean, median or most frequent (mode) value.

- Pros: use the whole dataset
- Cons: reduce variance and the correlation between variables.

		Mean		Median		Mode	
Gender	Age	Gender	Age	Gender	Age	Gender	Age
F	20	F	20	F	20	F	20
M	NaN	M	31	M	30	M	40
F	25	F	25	F	25	F	25
M	30	M	30	M	30	M	30
F	40	F	40	F	40	F	40
F	40	F	40	F	40	F	40

```
from sklearn.impute import SimpleImputer  
imp = SimpleImputer(strategy = 'mean')
```

```
SimpleImputer(strategy = 'median')
```

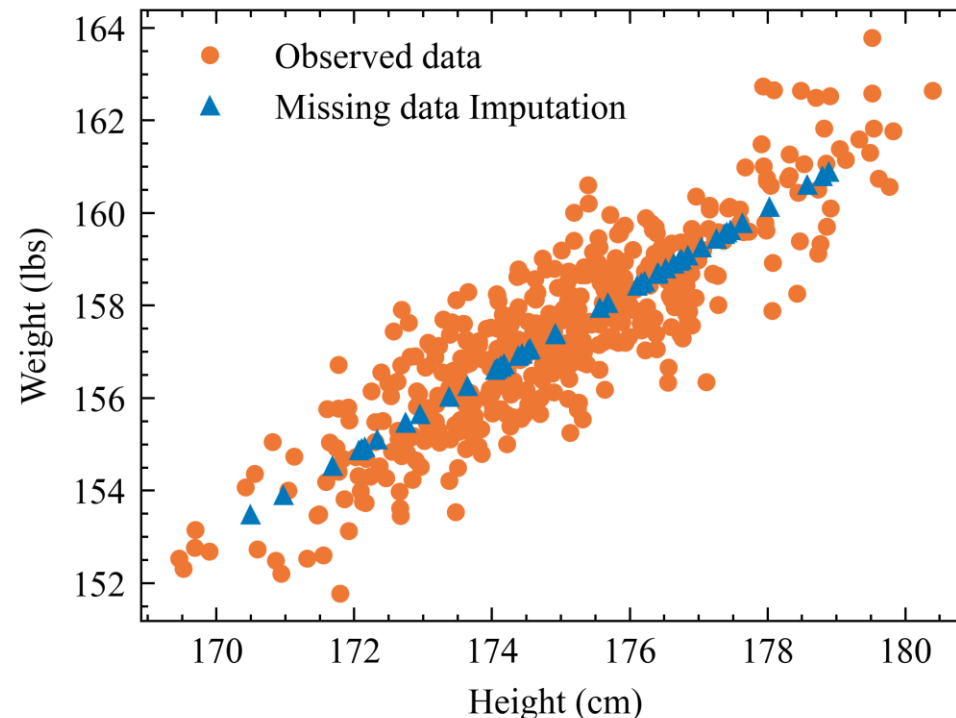
```
SimpleImputer(strategy = 'most_frequent')
```

Handling Strategies – Machine Learning Model Inference

Predictive/Statistical models to infer the values of missing data. There are many options for such predictive model – Linear regression / Random Forest / KNN / Neural Networks...

- Pros: use information from the observed data; can be effective with cross-validation.
- Cons: over-estimate correlation.
- Linear Regression

Height (cm)	Weight (lbs)
170	150
171	NaN
175	155
180	162
177	NaN
178	160



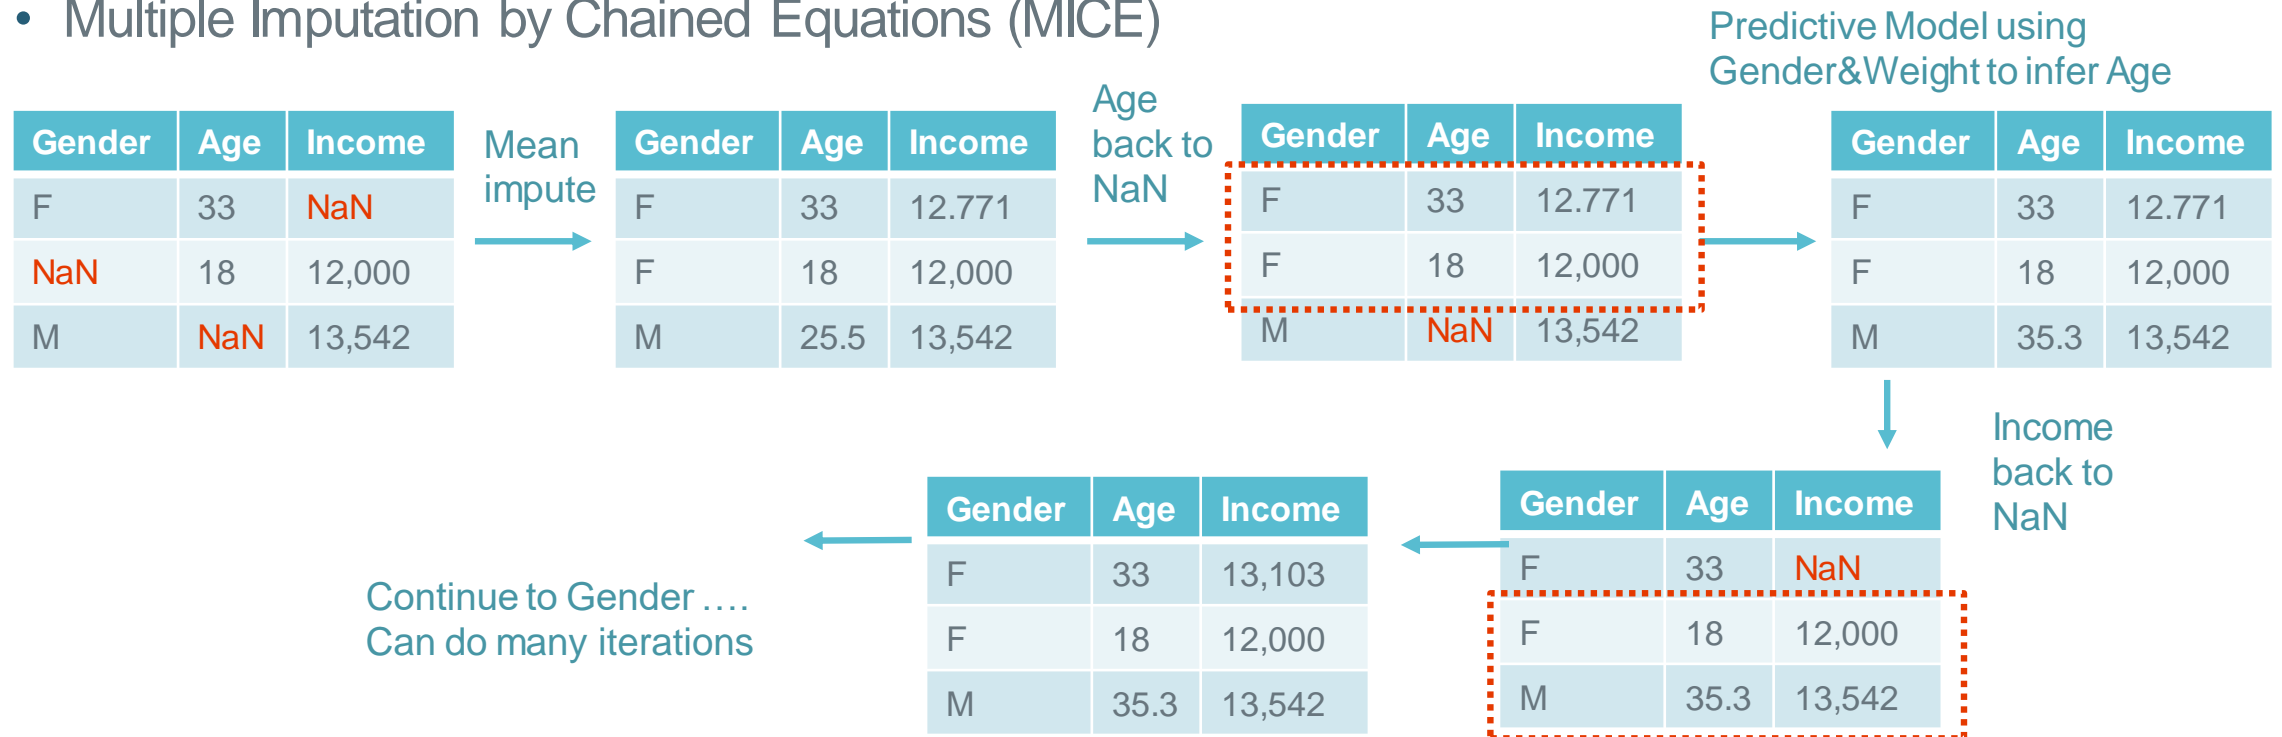
- Depends on the model performance
- Follow the assumed relationship of the model

Handling Strategies – Multiple Imputation

Multiple imputation: missing values are filled multiple times to create “complete” datasets.

- Cons: Having multiple values reduces bias.
- Pros: Highly technical and difficult to implement.

- Multiple Imputation by Chained Equations (MICE)



Python Examples

- [Notebook with these examples on Github](#)

Questions?

Future Topics

Have requests?
Let me know!

Introduction to Data Science Course Outline

Andrew Wheeler, PhD, andrew.wheeler@hms.com

- Lesson 01: Data Science 101
- Lesson 02: Machine Learning 101
- Lesson 03: Evaluating Predictions
- Lesson 04: Intro Data Transformation in Python
- Lesson 05: Data Visualization 101
- Lesson 06: Feature Engineering
- Lesson 07: Missing Data
- Lesson 08: Big Data and Parallel Computing Intro
- Lesson 09: Dimension Reduction and Unsupervised Learning
- Lesson 10: High Cardinality (Many Categories)
- Lesson 11: Intro to Forecasting
- Lesson 12: Conducting Experiments



Missing Data

Data Science & Machine Learning Team

03/31/2021

Jingjie Zhang, PhD

jingjie.zhang@hms.com